

UiT INF-3200 Distributed Systems - Project 1

Fall 2015

Enrico Tedeschi
ete011@post.uit.no

Mike Murphy
mmu019@post.uit.no

I. INTRODUCTION

Our task was to implement a simple distributed key-value store.

The general idea is to have a number of back-end nodes that store data, and a front-end node that forwards requests into the back-end group. The front-end should be able to contact any storage node and get the same results.

A. Requirements

Our data store ...

- must incorporate multiple storage nodes.
- must work without any node having complete knowledge of the others.
- must work no matter which storage node the front-end contacts. For testing, the front-end should contact random nodes.
- does not need to support dynamic adding and removal of storage nodes.
- does not need to persist data between runs.

II. TECHNICAL BACKGROUND

The basic goal of a distributed database is to coordinate computers so that they function as if they were a single machine. In a peer-to-peer system with no designated master node, there must be a deterministic algorithm, shared by all nodes, that decides which node or nodes should store which values. Each node must be able to determine whether or not it should be responsible for a given bit of data, and if not, it must be able to forward the request to the node that is. With a large number of nodes, it becomes impractical for all nodes to keep a list of all of the others, so each node will only keep a list of a small number of neighbors, and the distributed algorithm will have to route requests through this network of neighbor nodes to the correct location. [2]

A key-value store lends itself well to this kind of peer-to-peer system, because the set of possible keys can be viewed as a space to be divided into regions to

be handled by different nodes. The whole system then functions as a distributed hash table, or DHT. The key questions in designing a DHT system are how the key space should be divided among nodes, and which nodes should be neighbors and communicate. [2]

III. DESIGN

In our project, we arranged our key space along one dimension in a ring, following the lead of Chord[3]. Regions of the key space are assigned to each node. Each node is aware of the range of key space that it is responsible for, and also the next node in the ring (Fig 1). If a node receives a request to store or retrieve a key that is not in its range, it will forward the request to the next node.

Though our design was inspired by Chord, we did not implement Chord's finger table optimization, for lack of time.

Because all nodes are known ahead of time, with no joining or leaving, we could divide the key space evenly between each node.

For forwarding requests to successor nodes, we used a simple synchronous strategy: the node's request-handler thread will block while it contacts the next node. This was the easiest to implement, but it is wasteful of resources.




Fig. 1: Design of the implemented Network

IV. IMPLEMENTATION

A. Languages and Code

Our solution is implemented in a mix of Python and Bash script, Python for the actual node and front-end programs, and a Bash script to start up and shut down nodes via SSH.

We started with skeleton code by our teaching assistants, Einar Holsbø Jakobsen and Magnus Stenhaus, which included the front-end code and the startup/shutdown shell script. Their front-end also included an automated test routine that rapidly inserted and retrieved random key-value pairs.

We wrote the actual storage node code, and we also made a few enhancements to the startup shell script. We modified the startup script to accept its list of nodes from the command line, rather than a list inside the script itself. We also added another short script to pick random hosts from the cluster and print them as a list ready to be copied and pasted into the command line. This made it easier to randomize which hosts in the cluster we were using, in hopes of minimizing conflict with other users on the cluster.

B. Network Protocol

The backend nodes accept and retrieve data through a simple HTTP API. HTTP's PUT and GET operations are a natural fit for a key-value store, as their semantics specify storing and retrieving documents (value) at a given URL path (key). Jakobsen and Stenhaus chose this protocol for their starter front-end node and we reused it for the storage nodes.

C. Persistence

The purpose of this exercise was to investigate the challenges of distributed data storage, not storage itself. Therefore, there was no requirement to actually persist stored data between runs. So, for simplicity, we did not implement any kind of data persistence. Data is simply stored in memory and the store starts empty on each test run.

D. Frontend

As specified in the requirements, the *Frontend Node* (Fig 2) contacts a random backend node for each request, and the backend nodes must cooperate to handle the request. This randomness enforces the distribution transparency of the data store. The system can store and retrieve keys as a whole, no matter which individual node is contacted.

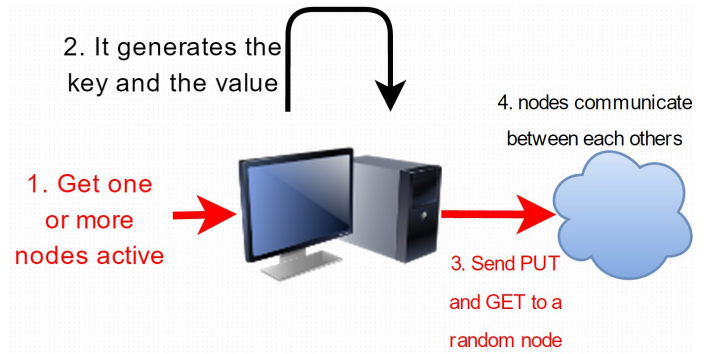


Fig. 2: Implementation of the Frontend node

E. Node

Each node (Fig 4) has a simple workflow. It starts running and waits for a request, which could be GET or PUT from the frontend node or another backend node. For each request, the node hashes the key into the linear key space, and checks if it falls within the region assigned to that node. If so, the node handles the request by storing the value (for PUT) or returning a previously stored value (for GET). If the key does not fall in the node's assigned range, it forwards the request to the next node (Fig 3).

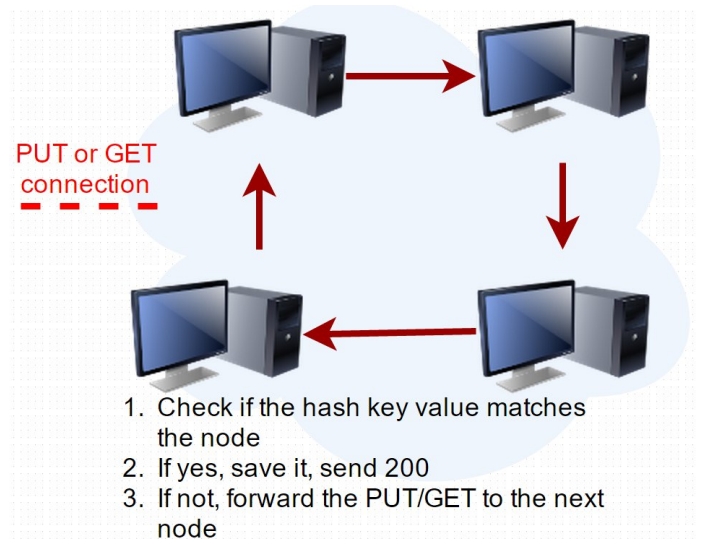


Fig. 3: Implementation of the PUT/GET calls

The node's logic is split across two classes: *NodeCore* and *NodeHttpHandler*. *NodeCore* includes the core logic of deciding when to store a key and when to forward a request to the next node. *NodeHttpHandler* includes the logic of interpreting HTTP requests and formatting HTTP responses. This separation of logic makes it

possible to verify the core algorithm with isolated unit tests, without having to set up HTTP servers.

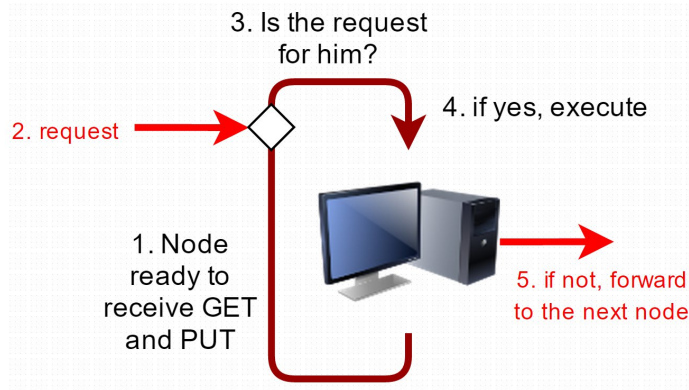


Fig. 4: Implementation of Individual Node

For hashing, the node uses the MD5 algorithm to map the string key to a numeric value, then takes that large integer modulo the number of nodes to get a node number (Fig 5). As a cryptographic algorithm, it is deterministic and it will evenly distribute the keys. MD5 is also fast[4], so it will not slow down lookups unnecessarily. And though it no longer considered suitable for security applications[5], it is not being used for security here. We merely need to decide which node should store a given key.

Fig. 5: Pseudocode of Core Node Logic

```

1 node_responsible = MD5(key) % node_count
2 if node_responsible == self.rank
3 then
4     save or retrieve value
5 else
6     forward request to next node
  
```

F. Environment

Our code was written to run on the Rocks Cluster distribution[1], and makes some assumptions about that environment. We rely on the cluster's shared filesystem for distributing program code to servers. And we rely on easy SSH access between machines in the cluster to start and shutdown nodes.

V. DISCUSSION

Our design decisions favored simplicity over performance. The simple ring structure was easy to implement, but requests may need to be forwarded along several nodes in the ring before finding the correct key. The number of hops is linear with the number of nodes,

$O(n)$. Chord's finger-table optimization[3] finds nodes in $O(\log_2 n)$ hops. We had hoped to implement the same strategy in our project but we ran out of time.

The simple synchronous request forwarding strategy is also a potential bottleneck. It leaves a thread idle on each node involved in each request, waiting until the right key is found before returning it back through all of the nodes that were involved. We suspect that, as the number nodes grows or request frequency increases, this holding of resources will choke the system. The advantage of this approach is its simplicity. The request and response protocol is identical between client and front-end, front-end and storage node, and between storage nodes themselves. The storage nodes also use the same logic to handle requests whether the request comes from the front-end or another node in the chain.

Asynchronous message-passing would allow intermediate nodes in a search to free resources. The first node could block, and then send a message to the next node. If this message included a return address to the first node, then the other nodes in the ring could pass it along without leaving connections open or blocking. Finally, the node that has the key could send the value to original node. When that original node received its answer, it could unblock and send the response to the front-end. Such a strategy should be more efficient, but would require different message formats for requests and answers, and nodes would need logic to receive answer messages and match them to waiting threads to send responses. We suspect that the finger tables would give a much larger performance boost. Not only that, but with $O(\log_2 n)$ searching, only a few nodes would be involved in each search, and the overhead from waiting synchronously on just a few nodes might be acceptable.

VI. EVALUATION

For the evaluation node scaling has been considered. The function *storage_frontend* has been timed sending five hundred requests (GET/PUT) to the nodes network. The evaluation has been done considering the scale on the number of nodes and for each, 10 tests were taken into consideration.

The number of nodes and the average time of 10 computations is represented in the table in Fig 6.

In the Fig 7 instead the graphic of this scaling test is characterized.

VII. CONCLUSION

Our DHT solution, with a simple ring structure, was able to store and retrieve data correctly, in time that increased linearly with the number of nodes ($O(n)$).

Fig. 6: Nodes/Time scaling table

Nodes	Time
2	5.7923
4	6.4793
6	8.1309
10	13.0746
15	17.0469
20	19.3472
30	29.4729
40	35.2886

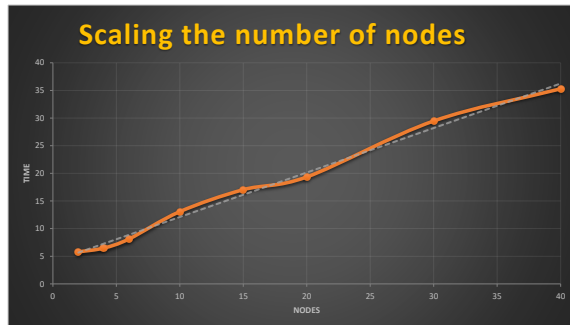


Fig. 7: Scaling considering the number of nodes in the network

REFERENCES

- [1] Rocks clusters: About. http://www.rocksclusters.org/wordpress/?page_id=57.
- [2] Hari Balakrishnan, M. Frans Kaashoek, David Karger, Robert Morris, and Ion Stoica. Looking up data in P2P systems. *Commun. ACM*, 46(2):43–48, February 2003. <http://doi.acm.org/10.1145/606272.606299>.
- [3] Frank Dabek, Emma Brunskill, M. Frans Kaashoek, David Karger, Robert Morris, Ion Stoica, and Hari Balakrishnan. Building peer-to-peer systems with chord, a distributed lookup service. <https://pdos.csail.mit.edu/papers/chord:hotos01/hotos8.pdf>. MIT Laboratory for Computer Science.
- [4] Ronald Rivest. The MD5 message-digest algorithm. RFC 1321, RFC Editor, April 1992. <http://dx.doi.org/10.17487/RFC1321>.
- [5] Sean Turner and Lily Chen. Updated security considerations for the MD5 message digest and the HMAC-MD5 algorithms. RFC 6151, RFC Editor, March 2011. <http://dx.doi.org/10.17487/RFC6151>.